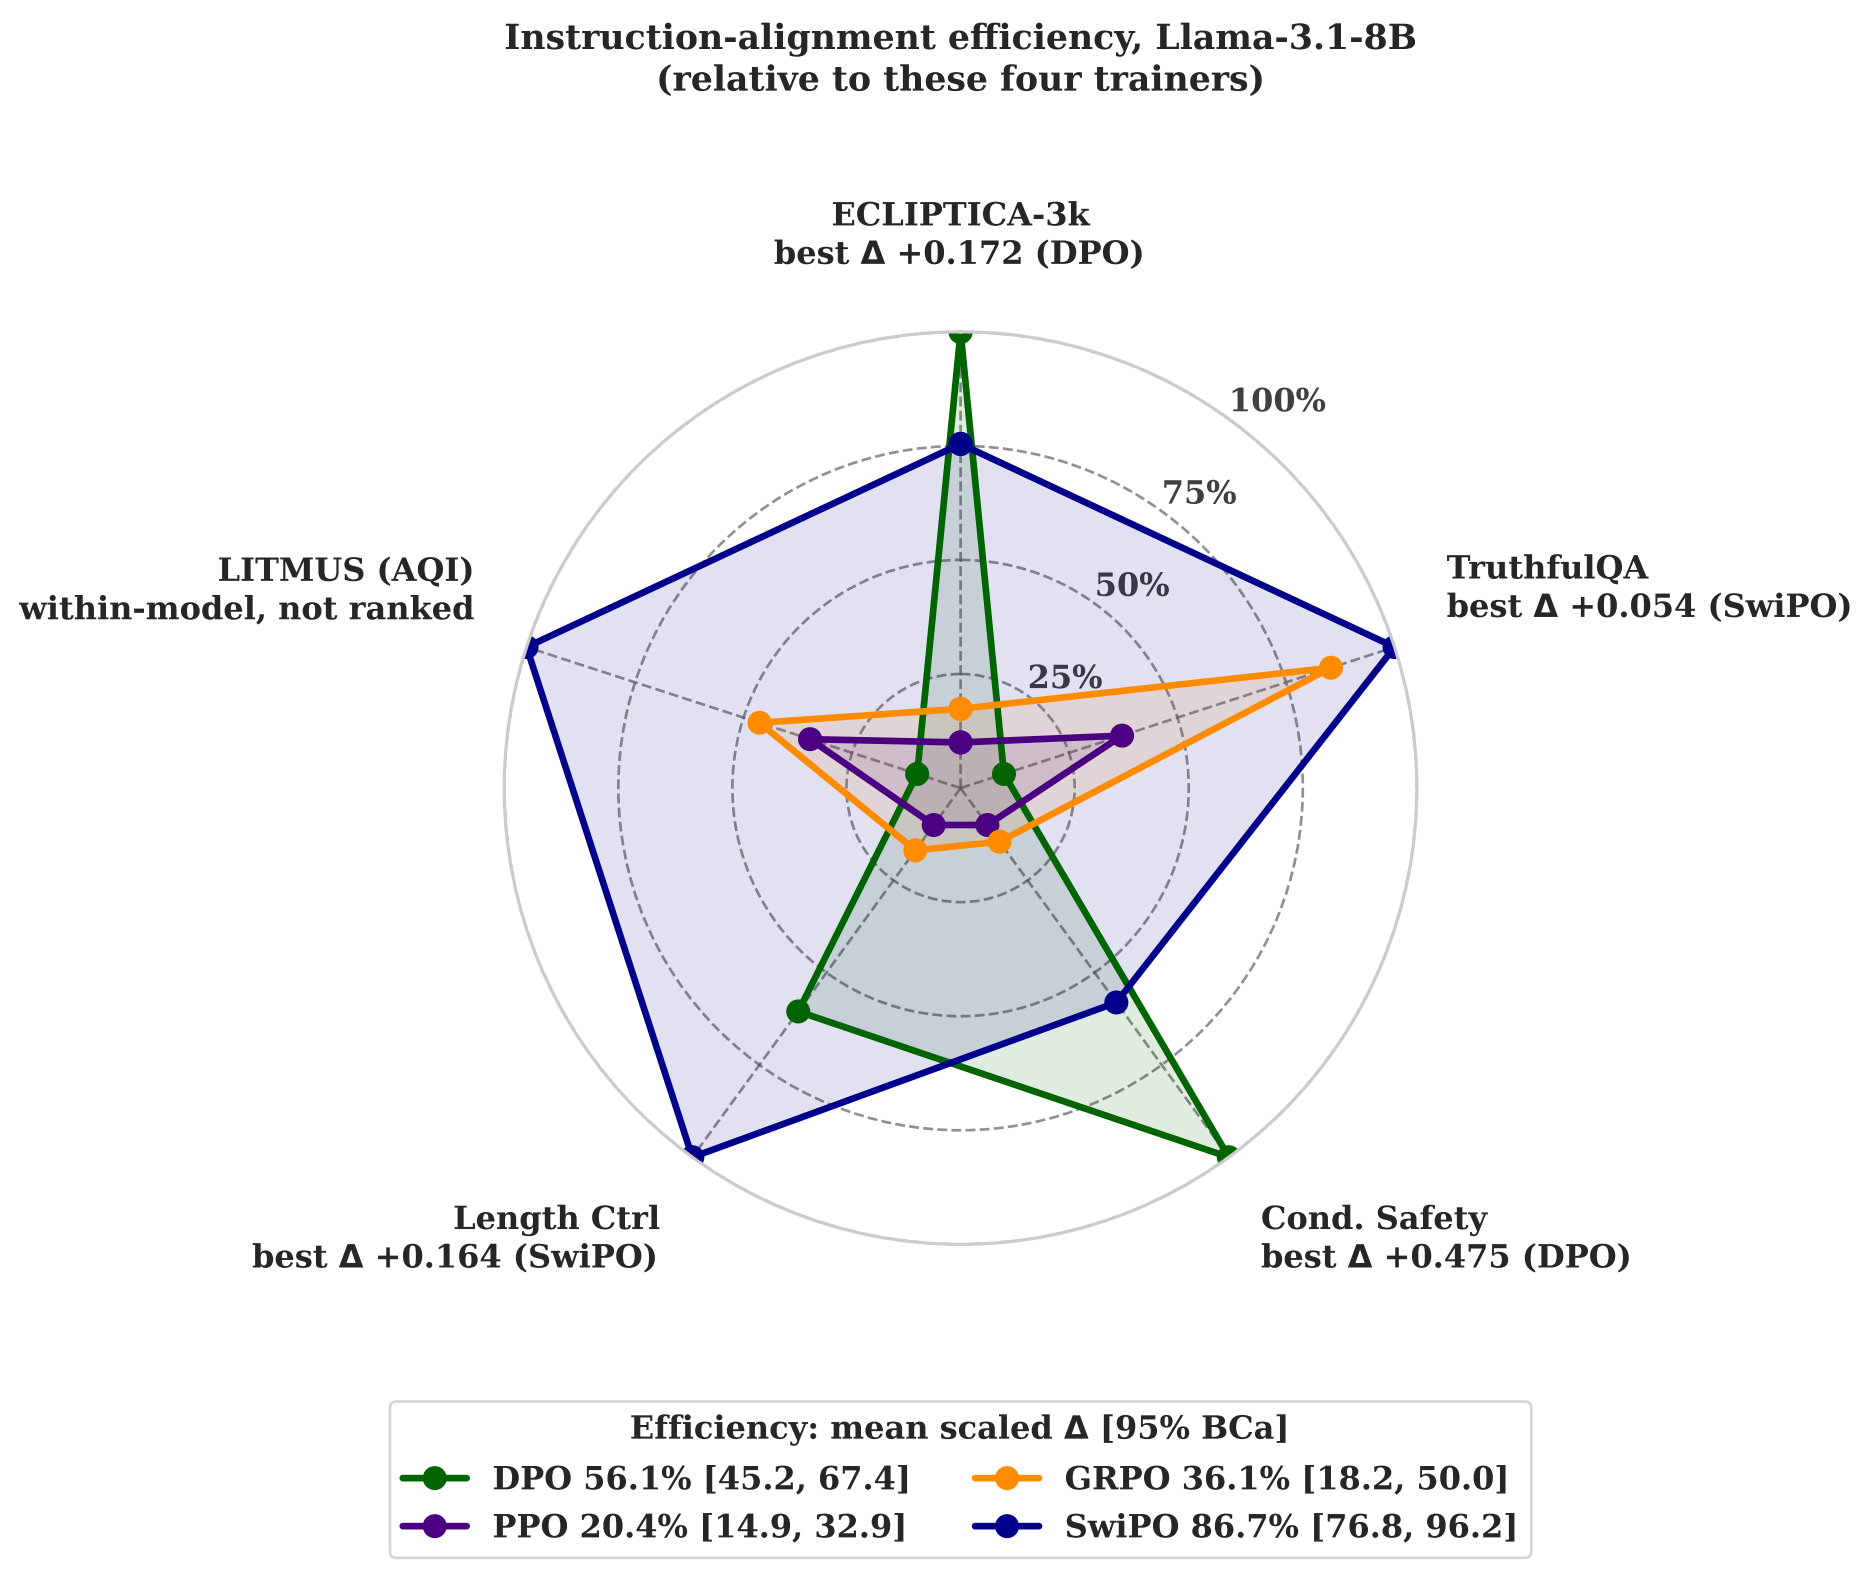
Instruction-alignment efficiency, Llama-3.1-8B
(relative to these four trainers)
ECLIPTICA-3k
best Δ +0.172 (DPO)
100%
75%
50%
25%
TruthfulQA
best Δ +0.054 (SwiPO)
Cond. Safety
best Δ +0.475 (DPO)
Length Ctrl
best Δ +0.164 (SwiPO)
LITMUS (AQI)
within-model, not ranked
Efficiency: mean scaled Δ [95% BCa]
DPO 56.1% [45.2, 67.4]
GRPO 36.1% [18.2, 50.0]
PPO 20.4% [14.9, 32.9]
SwiPO 86.7% [76.8, 96.2]
Vertices: Δ min-max scaled to [0.1, 1] per suite (no interval);
raw Δ with 95% intervals in the score table