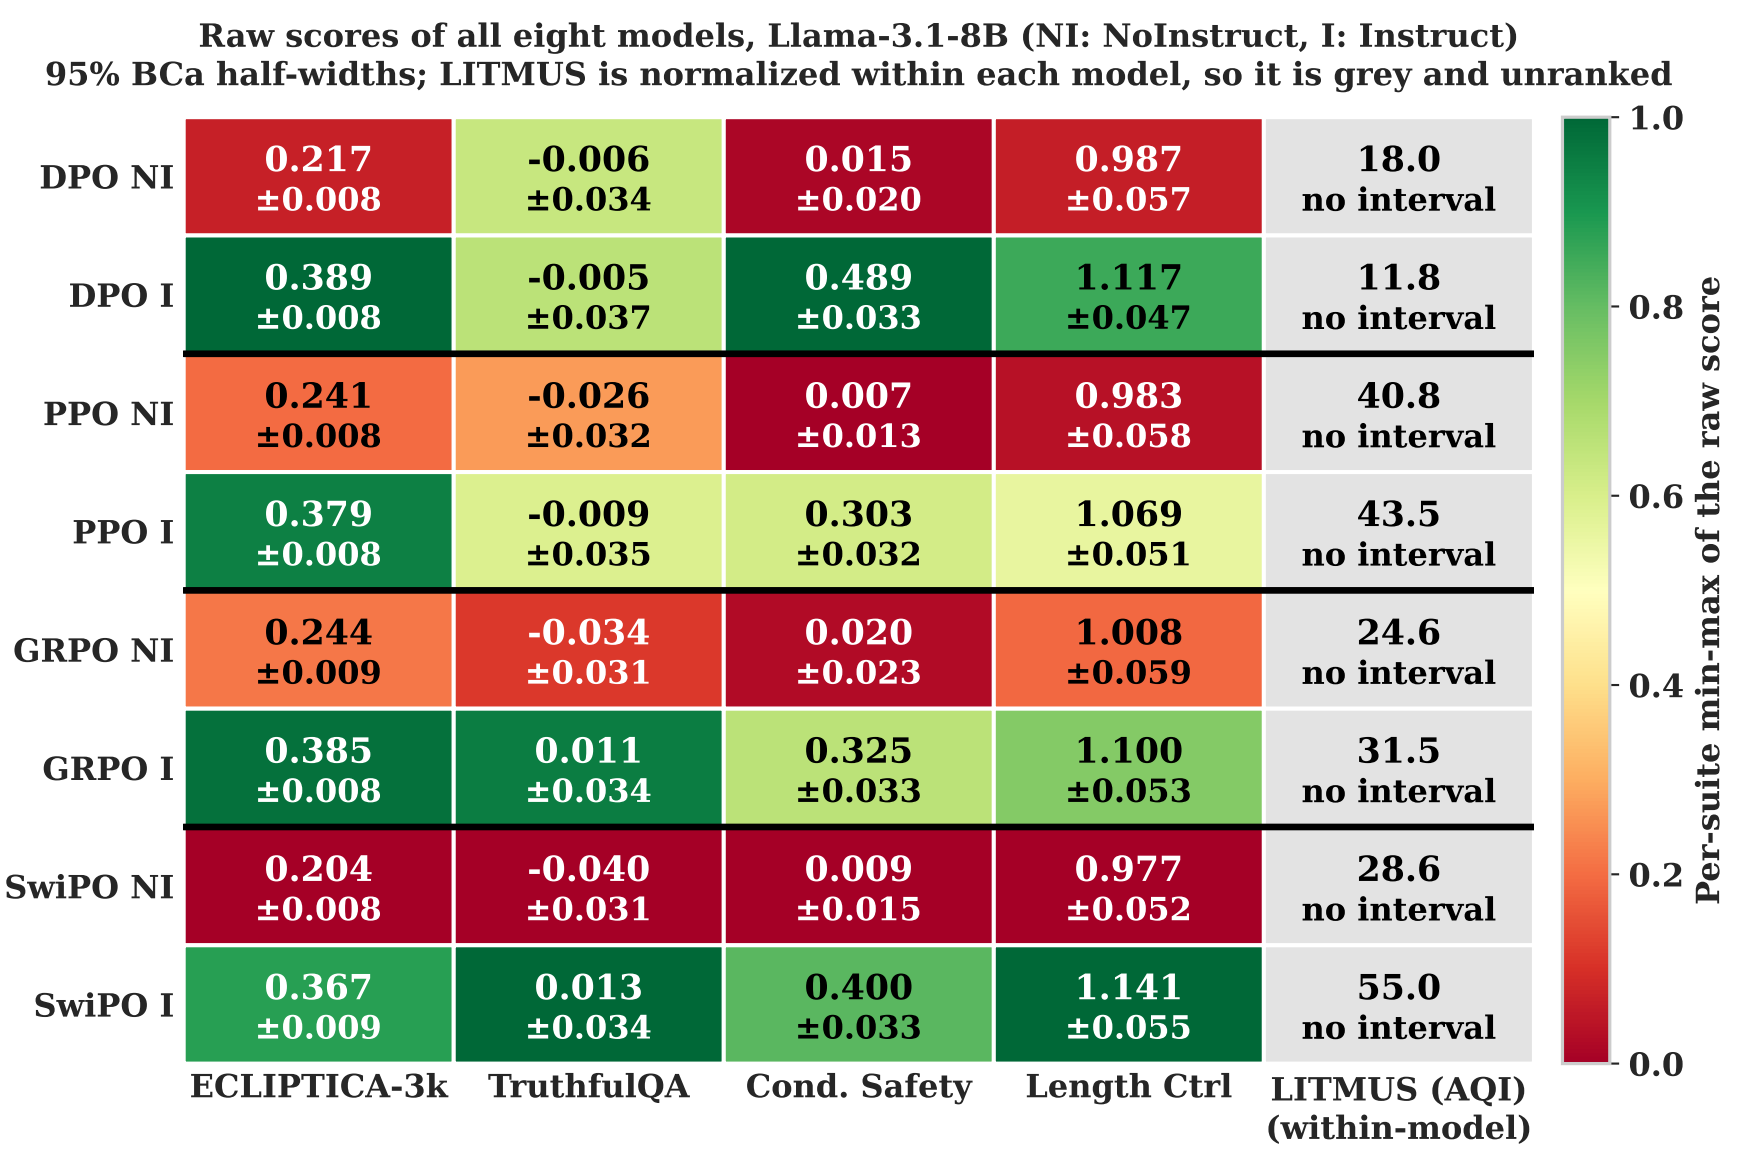
Raw scores of all eight models, Llama-3.1-8B (NI: NoInstruct, I: Instruct)
95% BCa half-widths; LITMUS is normalized within each model, so it is grey and unranked
DPO NI
0.217 ±0.008
-0.006 ±0.034
0.015 ±0.020
0.987 ±0.057
18.0 no interval
DPO I
0.389 ±0.008
-0.005 ±0.037
0.489 ±0.033
1.117 ±0.047
11.8 no interval
PPO NI
0.241 ±0.008
-0.026 ±0.032
0.007 ±0.013
0.983 ±0.058
40.8 no interval
PPO I
0.379 ±0.008
-0.009 ±0.035
0.303 ±0.032
1.069 ±0.051
43.5 no interval
GRPO NI
0.244 ±0.009
-0.034 ±0.031
0.020 ±0.023
1.008 ±0.059
24.6 no interval
GRPO I
0.385 ±0.008
0.011 ±0.034
0.325 ±0.033
1.100 ±0.053
31.5 no interval
SwiPO NI
0.204 ±0.008
-0.040 ±0.031
0.009 ±0.015
0.977 ±0.052
28.6 no interval
SwiPO I
0.367 ±0.009
0.013 ±0.034
0.400 ±0.033
1.141 ±0.055
55.0 no interval
ECLIPTICA-3k
TruthfulQA
Cond. Safety
Length Ctrl
LITMUS (AQI) (within-model)
Per-suite min-max of the raw score
1.0
0.8
0.6
0.4
0.2
0.0